

An Analysis on Chronic Kidney Disease Public Health Data

Akshay Sanjeev

July 11, 2025

Preface

0.1 Notation and Convention

By standard dataset, from now on I will be referring to the following data structure. We have m independent variables(a.k.a features), denoted by $x^i, i = 1, 2, 3, \dots, m$ and one dependent variable(a.k.a label), m . Let n be the total number of observations, and the lower index denotes the observation. Thus the dataset is set like the one given below.

$$\begin{pmatrix} x_1^1 & x_1^2 & \cdots & x_1^m & y_1 \\ x_2^1 & x_2^2 & \cdots & x_2^m & y_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ x_n^1 & x_n^2 & \cdots & x_n^m & y_n \end{pmatrix} \quad (1)$$

0.2 Model Evaluation

| | Predicted Positive | Predicted Negative |
|------------------------|---------------------------|---------------------------|
| Actual Positive | TP (True Positive) | FN (False Negative) |
| Actual Negative | FP (False Positive) | TN (True Negative) |

Accuracy

Accuracy is the proportion of all predictions (both positive and negative) that were correct. It is not recommended when dealing with imbalanced datasets.

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

Recall(TPR)

Also known as the true positive rate or sensitivity, recall measures the proportion of actual positive cases that were correctly identified.

$$\text{Recall} = \frac{TP}{TP + FN}$$

False Positive Rate (FPR)

The false positive rate measures the proportion of actual negatives that were incorrectly classified as positive. Ideally, this should be close to zero.

$$\text{FPR} = \frac{FP}{FP + TN}$$

Precision

Precision is the proportion of predicted positive cases that are truly positive.

$$\text{Precision} = \frac{TP}{TP + FP}$$

0.2.1 ROC curve

ROC curve is plotted as true positive rate(TPR) on y axis and false positive rate(FPR) on x axis. Ideally we want TPR to be 1 and FPR to be 0. Hence the closer it follows the left and upper axis, better the results. Another related measure is the area under the ROC curve. The best result would be $\text{AuC} = 1$.

1 Pre-Processing Data and Question

The motivation to pre- process data is that you need the dataset in a machine readable format. The raw dataset, titled `CKD_agri_data.xls` is in the following format. 3350 rows correspond to individuals and 151 columns contain various information about these individuals. The columns have various information It has metadata like name, age, sex and so on, medical information, disease history, test reports and etc. This is a mix of numeric and categoric variables.

The first issue I found was the missing observations. I looked at the what fraction of data is missing for each row and column and plotted the sorted fraction in figure 1. I set thresholds for the both row-wise and column-wise fraction and only kept the observations below the threshold. The blue line in figure 1 represents the threshold in both cases.

Second issue was that some features had erroneous feature names and values. I corrected the feature names and replaced the erroneous values with the median of the respective feature. At this point, the dataset was mostly clean and ready for further analysis. For categorical variables with erroneous or missing values, I replaced them with the mode (most frequent value) of the respective feature.

I had 46 numeric features and 54 categoric features now. To begin the analysis I only chose the numeric features and created the dataset. Among the numeric dataset, I had four features related to Chronic Kidney Disease(`ckd_final`, `ckd_epi_sample_1`,`ckd_probable_sample_1`, `ckd_code_sample_1`).

But since the only information I have is the feature name, it is hard to predict what they mean. But looking at the unique values in each column, I inferred that `ckd_final` and `ckd_probable_sample_1` are binary variables and `ckd_probable_sample_1` has 0, 1, 2 and 3 as unique values and the last one is roughly in the range 0-145. The histogram of these four variables are given in figure 2.




Figure 1: Fraction of missing data corresponding to each column(Top) and each row(Bottom).

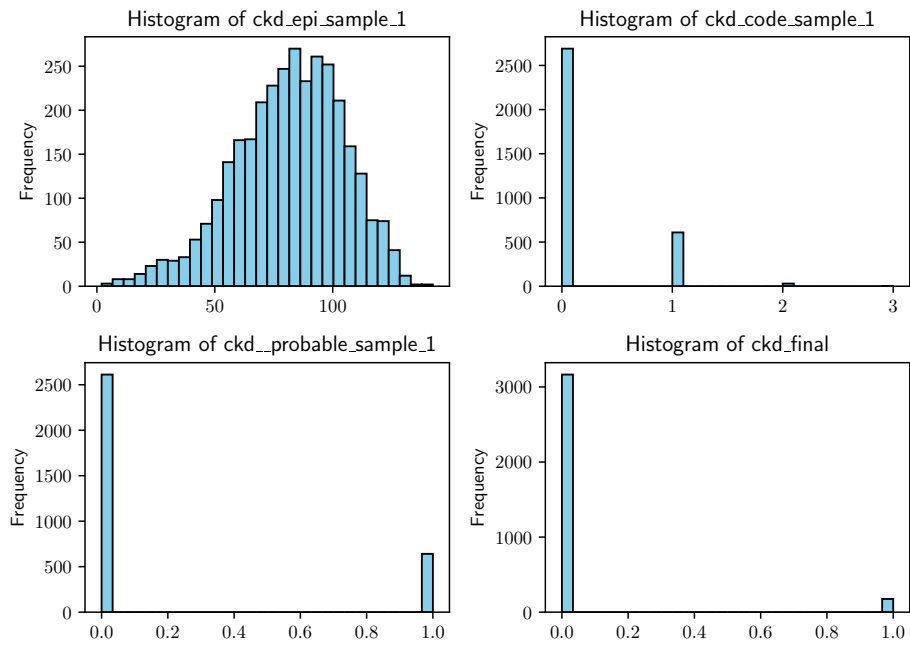


Figure 2: Distribution of potential target variables




Figure 3: Each subplot is trying to visualise relation between target variable, `ckd_code_sample_1` with other 3 variables related to chronic Kidney disease.

1.1 Question

Since the dataset primarily target on chronic Kidney disease, the first question we are intersted in is given the numeric data collected, can we use machine learning based methods to predict whether a person having chronic kidney disease or not. I tried to find a correlation between the numeric CKD feature and other features in pursuit of picking a target variable among the four. The plots are in figure 1.1. Qualitatively, `ckd_probable_sample_1` has the highest correlation. But I think it's okay to remove all four from training data to be safe. For the target variable, I choose `ckd_code_sample_1`. Presumably, the variable has the optimal information, because it predicts 3 stages of chronic Kidney disease.

Now that the question is clear, we will use Decision Trees, Random Forest and Feed-froward Neural Network models on the data.

2 Decision Tree

2.1 Informal Description(DT)

A decision tree is a supervised learning method, where you split data based on a nested `if` and `else` conditions. To build a decision tree,

1. take the standard dataset, compute the Shannon Entropy for each variable

$$H^i = - \sum_j^n p(x_j^i) \log p(x_j^i)$$

2. Classify or predict the dependent variable based on those each variable and compute the Information Gain(IG^i) for each variable,

$$IG^i = H^i - \sum_s \frac{|x^{is}|}{|x^i|} H^i_s$$

, where x^{is} is the set of s -th split of dataset.

3. Pick the feature with maximum IG and assign that for the node.

$$x_{\text{selected}} = \arg \max_i (IG^i)$$

Repeat the step one to three iteratively by moving down the tree, as you pick a feature and assign it to a node, until all nodes become terminal. A node is terminated if the it has an zero entropy.

2.1.1 Sensitivity

This is a greedy splitting. At each step the aim is to make immediate improvements, which may not be the best at a global level. It is a local phenomenon. Also if we don't set a maximum depth, trees grow deeper and may end up in memorizing the training data and over-fitting. This overall results in poor generalisation and high variance. This can be tackled with a couple of methods (pruning, specified tree depth and minimum samples per terminal node, ensemble of trees etc.)

2.2 Results

I build a decision tree and carried out the analysis using Python 3.11 and scikit-learn a machine learning package in Python. The tree was built using the Entropy criterion and the rest of the parameters are kept default. The classification results are represented as a confusion matrix in figure 2.2. The classifier did a very good job for class 0 and 1. But possibly due to the scarcity in training data, the the class 2 didn't perform well. Class 3 was too rare, that it was not present in testing data.

The contribution of each feature in figure 2.2. The highest contributing features are **dbp2** and **gft_cat_sample_1**. To conclude from public health data we can infer potential chronic Kidney disease patients, given it contains reliable medical test information, as we have seen above test results like **dbp2** and **gft_cat_sample_1** are the most important features. Qualitatively all other features seems to contribute very minimally. But some statistical significance measure is recommended for a reliable conclusion.




Figure 4: (Left) Confusion matrix of Decision tree classification. Precisions values are 0.98, 0.83 and 0 for class 0, 1 and 2 respectively. Accuracy is 0.94, (Right) The 10 highest contributing features to the model.

3 Random Forest(RF)

3.1 Informal Description

Random Forest is an ensemble of Decision Trees. The justification for building an ensemble of Decision Trees is to reduce variance and sensitivity. To be specific, it uses a method called Bootstrap Aggregation. First sample various sets of observations and features with replacement from the independent decision trees. The final prediction is done by using averaging or major-voting depending on the varibale. For example we have a model, $\hat{f}^i(x)$, ($i = 1, 2, 3, \dots, B$) and B bootstrapped datasets and our final prediction for a regression will be

$$\hat{f}_{\text{avg}}(x) = \frac{1}{B} \sum_{i=1}^B \hat{f}^i(x)$$

3.2 Analysis & Results

Similar to the above case, I have used Python 3.11 and scikit-learn to build Random Forest, classify and analyse the results. Rather jumping into the classification problem we have, at first I was interested in is how would the algorithm perform as you increase the number of trees. So I computed the accuracy of random forest, with 1 to 600 trees (figure 3.2). The accuracy of classification saturates for both training and validation dataset a little early than 100 trees. Thus I picked 100 trees and did the classification.

The summary of a result with confusion matrix and feature importance are plotted in 3.2. As we have seen for Decision tree, class 0 and 1 are fairly good. Now class 2 also doesn't have any




Figure 5:

observations in testing data. We can see a much wider feature importance here. Interestingly the highest contributing feature earlier is shifted to third place. The importance of leading features have also lowered, implying a more balanced classification.

From the ROC curve in the figure 3.2 also implies a fair classification for class 0 and 1. Even though no class 2 observations were predicted correctly, class 2 having a high area under the ROC curve(0.9284) is hard to understand.

To recapitulate, switching from decision tree to random forest, had given us marginal increase in accuracy. But I think the ensemble effect can be seen when we have large data to test the model on.

4 Feed-forward Neural Network(FFNL)

4.1 Informal Description

Feed-forward Neural Network is a supervised machine learning algorithm, based on a network of neurons or perceptron to be specific. Each neuron behaves like a switch, depending on the inputs signals the neuron decides to turn ON or OFF, by plugging the inputs to an activation function. The network architecture is layered as shown in the figure 7. Information flows from layer to layer, but not between nodes in the same layer.

For the i -th neuron on l -th layer, you get information from $l - 1$ th layer and your input to this neuron $z^{(l)}_i$ is,

$$z^{(l)}_i = W^{(l)}_{ij} a^{(l-1)}_j + b^{(l)}_i$$

where $W^{(l)}_{ij}$ is the weight of the matrix that goes from j to i , $a^{(l-1)}_j$ is the output of j -th neuron in the previous layer. This input is plugged into an activation function, to compute the output of the neuron.

$$a^{(l)}_i = f^{(l)}_i(z^{(l)}_i)$$

. The popular choices for activation function are ReLU, sigmoidal, Softmax, etc., We initiate and train the network with small random weights and make a prediction. Then compute how far is the

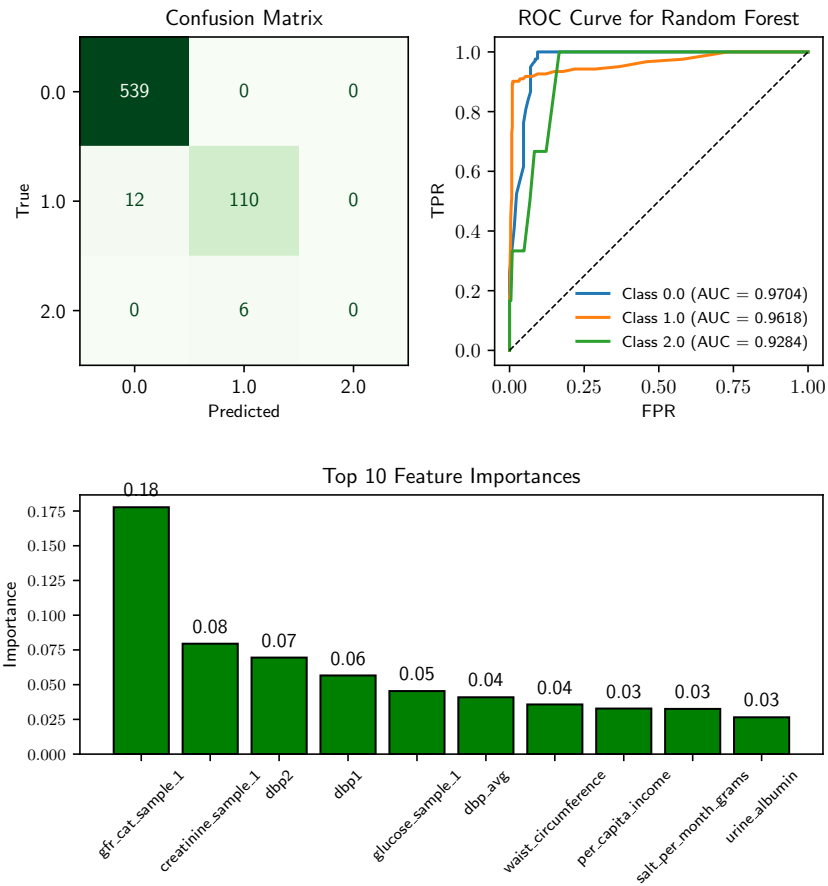


Figure 6: (Top, Left) Confusion matrix of Random Forest classification. Precisions values are 0.97, 0.95 and 0 for class 0, 1 and 2 respectively. Accuracy is 0.97, (Top, Right) ROC curve for different classes. (Bottom) The 10 highest contributing features to the model.




Figure 7: Schematic for a neural network

prediction from the desired output, using a cost function. The next step is to minimize the cost, by tuning the set of parameters, θ . You tune the parameters along the negative cost gradient with a learning rate, ϵ . The parameter set for the next prediction is,

$$\theta_{t+1} = \theta_t - \epsilon \frac{\partial}{\partial p} C(y, \hat{y})$$

The cost gradient is estimated using chain rule. For the weight matrix, it can be written as:

$$\frac{\partial C}{\partial W_{ij}} = \frac{\partial C}{\partial a_i^L} \cdot \frac{\partial a_i^L}{\partial z_i^L} \cdot \frac{\partial z_i^L}{\partial W_{ij}} \quad (2)$$

$$= \frac{\partial C}{\partial a_i^L} \cdot f^{(l)'}(z_i^{(l)}) \cdot a_j^{(l-1)} \quad (3)$$

You iteratively do this until you reach the minimum cost possible. The prediction part where you move from input to output is called forward propagation and the parameter tuning part is called back propagation.

4.2 Analysis & Results

Using TensorFlow 2 in Python 3.11. At first, I built an FFNL with 1 input layer with 46 input layers (same as the number of features), one hidden layer with 32 neurons and the output layer with 4 neurons, because I have 4 class labels. For the first two layers, I have a sigmoid activation function and for the final layer I have softmax. I am training the network for 50 epochs. I have tried longer epochs, but the accuracy is saturated around 0.8, which doesn't look that bad. But




Figure 8: Confusion matrix for Feed-forward Neural Network. Precision of class 0, 1 and 2 are 0.81, 0 and 0 respectively.

the classification matrix had the real horror. In figure you can see that everything was predicted as class 0 and since all class 0, which dominated the training dataset, has been classified correctly the resulting accuracy was 0.81. I made a bigger network, with (45, 128, 64, 32, 4) and still the same. I varied the learning rate, ϵ and nothing changed. An interesting observation was the network is not learning. The `loss`, `va_accuracy` and `val_loss` are not changing at all and `accuracy` is fluctuating around 0.8.

To conclude, I couldn't optimise the Feedforward neural network model for this problem.